

# Ben Irastorza

[bmirastorza@gmail.com](mailto:bmirastorza@gmail.com) | [linkedin.com/in/ben-irastorza](https://linkedin.com/in/ben-irastorza) | [github.com/birastor22](https://github.com/birastor22) | US Citizen

## EDUCATION

### Georgia Institute of Technology

Dec 2026

*Bachelor of Science in Computer Science*

GPA: 3.6/4.0

- Concentration: Artificial Intelligence, Information Internetworks
- Relevant Coursework: Machine Learning, Systems and Networks, Database Systems, Computer Networking

## EXPERIENCE

### Incoming Software Engineer Intern

June 2026 – Aug 2026

*Capital One*

*McLean, VA*

- Selected for Capital One's competitive Technology Internship Program (Summer 2026).

### Data Science Intern

Aug 2025 – Dec 2025

*Southern Company*

*Birmingham, AL*

- Developed the organization's most accurate winter peak demand forecast by implementing an XGBoost model with tailored meteorological features and data center load-growth analysis, reducing MAPE by 47%.
- Reduced data retrieval time by ~97% by developing a Python CLI tool leveraging the AVEVA PI Web API, replacing a manual Excel-based ingestion process with automated shell scripts for monthly analyses.
- Engineered a real-time AVEVA PI Vision dashboard measuring data center load swings' impact on Area Control Error (ACE) deviations, giving management the evidence needed to pursue new operational regulations.

### Advanced Analytics Intern

May 2025 – Aug 2025

*UCB Biopharma*

*Smyrna, GA*

- Enabled identification of undiagnosed patients and risk stratification by building a BERT-based model that encodes longitudinal claims data into dense embeddings, supporting patient journey clustering across large populations.
- Strengthened risk assessment for an anti-seizure medication by developing a survival analysis model with left-truncation and right-censoring corrections, benchmarking internal outcomes against real-world evidence.
- Validated drug usage patterns for a rare pediatric epilepsy disorder by running subpopulation analyses on Komodo claims data in Snowflake, surfacing prescribing trends that informed commercial strategy.

### Software Engineer Intern

May 2024 – Aug 2024

*The Parker Avery Group*

*Atlanta, GA*

- Enabled faster executive decision-making by building a LangChain-based AI analytics agent that translates natural language questions into business metrics and automatically generates dashboards.
- Increased agent reliability by curating a human-validated question-to-SQL dataset and implementing Milvus semantic retrieval to surface similar queries, increasing accuracy by 16%.
- Scaled client-to-competitor product matching across thousands of SKUs by migrating to batched embeddings via the OpenAI Batch API, cutting embedding costs by 50%.

## PROJECTS

### Enhanced Spotify Wrapped | *Java, Android SDK, Firebase, GitHub Actions, Figma, Git*

- Built an Android app using the Spotify Web API that delivers personalized music recaps, tailored recommendations, and real-time AI-powered conversations with users' favorite artists.
- Automated build validation and deployment with a GitHub Actions CI/CD pipeline, reducing integration errors across a multi-sprint development cycle.
- Integrated Firebase Authentication with OAuth 2.0 to manage Spotify token exchange and session persistence, securing user credentials across sessions.

## TECHNICAL SKILLS

**Languages:** Python, Java, C, SQL, HTML, CSS

**Libraries/Frameworks:** LangChain, Scikit-learn, XGBoost, PyTorch, Transformers, NumPy, Pandas, Matplotlib

**Development Tools:** Docker, Git, Jira, Figma, Wireshark, Power BI

**Cloud/Databases:** Databricks, Snowflake, PostgreSQL, MySQL, Firebase, Milvus, AVEVA PI System